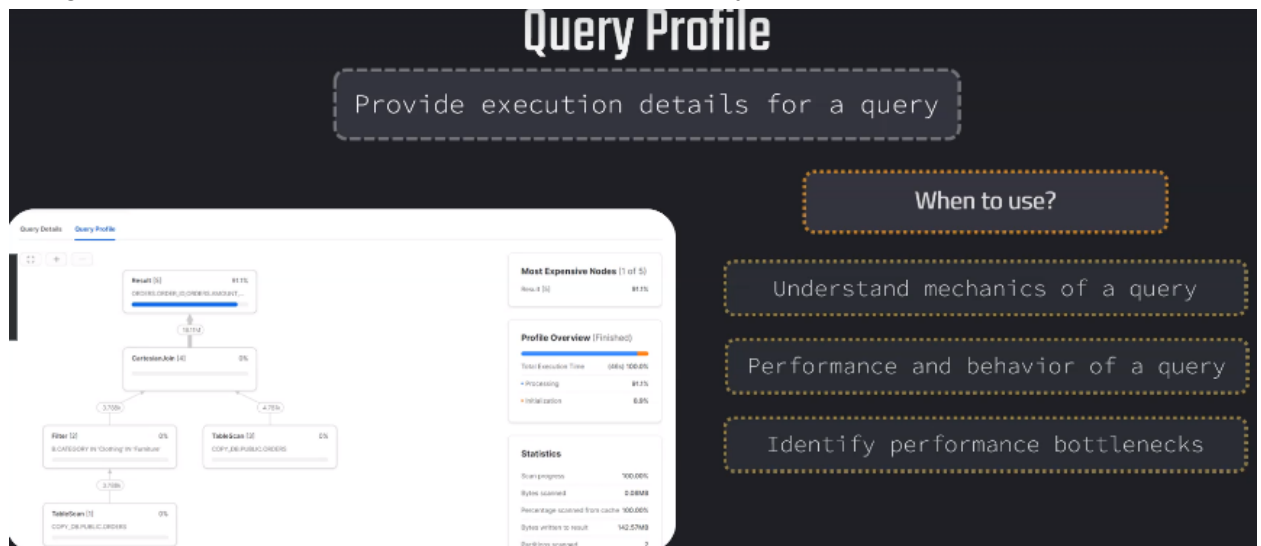


Query Profile & history


1. Query profile gives execution details
2. This is a graphical representation that we can use to get additional execution details for a given query.
3. Using this we can understand the mechanics of the query



- 4.
5. We can find the bottlenecks(point of lag) ..and can try to remove bottlenecks
6. It is available for all types of queries...like running,failed etc

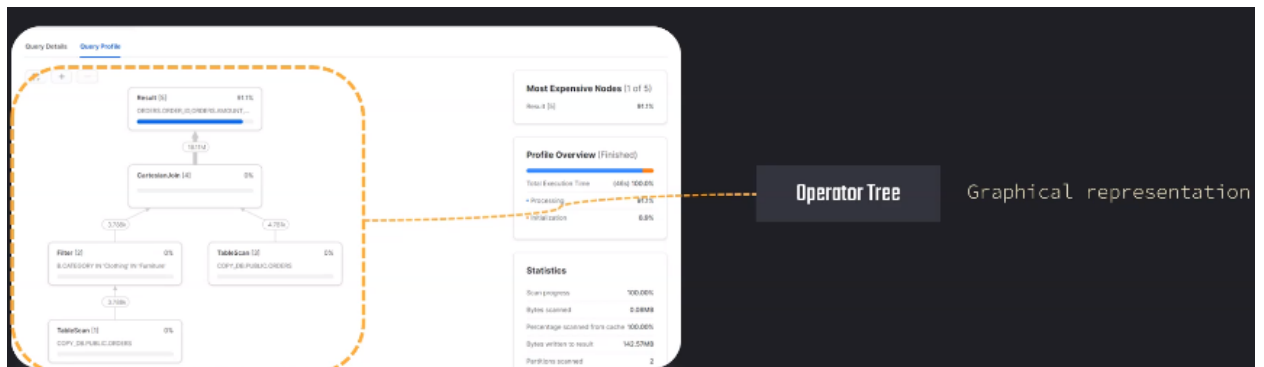
Available for all queries (Completed, Failed, Running)

SQL TEXT	QUERY ID	STATUS
SELECT * FROM SNOWFLAKE_SAMPLE_DATA.TPC...	01ab7153-0001-0b68-0003-a4c6000412...	Running
SELECT * FROM SNOWFLAKE_SAMPLE_DATA.TPC...	01ab7150-0001-0b68-0003-a4c6000412...	Success
SELECT * FROM SNOWFLAKE_SAMPLE_DATA.TPC...	01ab714a-0001-0b92-0003-a4c60003f1e6	Success
SELECT * FROM SNOWFLAKE_SAMPLE_DATA.TPC...	01ab713b-0001-0b92-0003-a4c60003f186	Failed

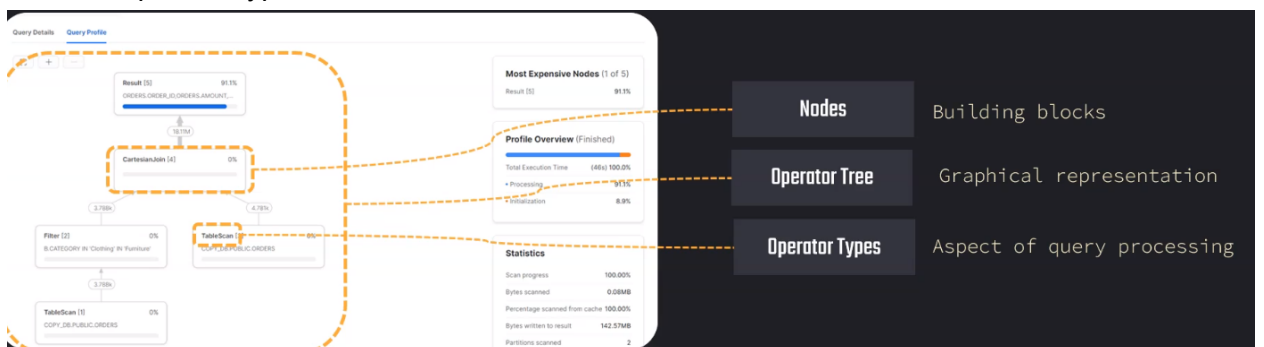


7. We can find query profile at


8. Now let's look into query profile...first we have an operator tree




9. In here we have different nodes...which are building blocks of our query...and each node has a different operator type



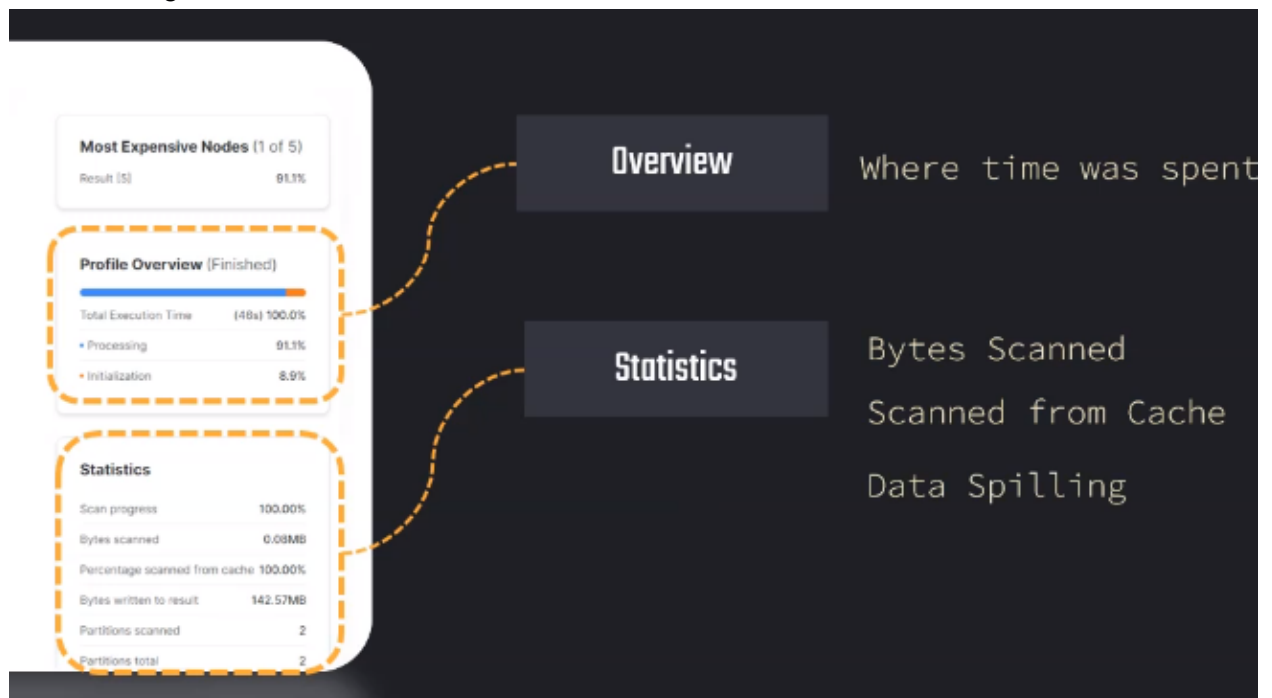
10. We can also see dataflow...it says that how many rows have been processed



11. Then we have a percentage



12. And to the right we have a dashboard



13. What is data spilling?

14. First we will see when does it occur

15. Lets suppose we have a small warehouse...when our data doesnt fit in this warehouse memory....it spills the memory to local storage....and in the end ..it decreases the performance

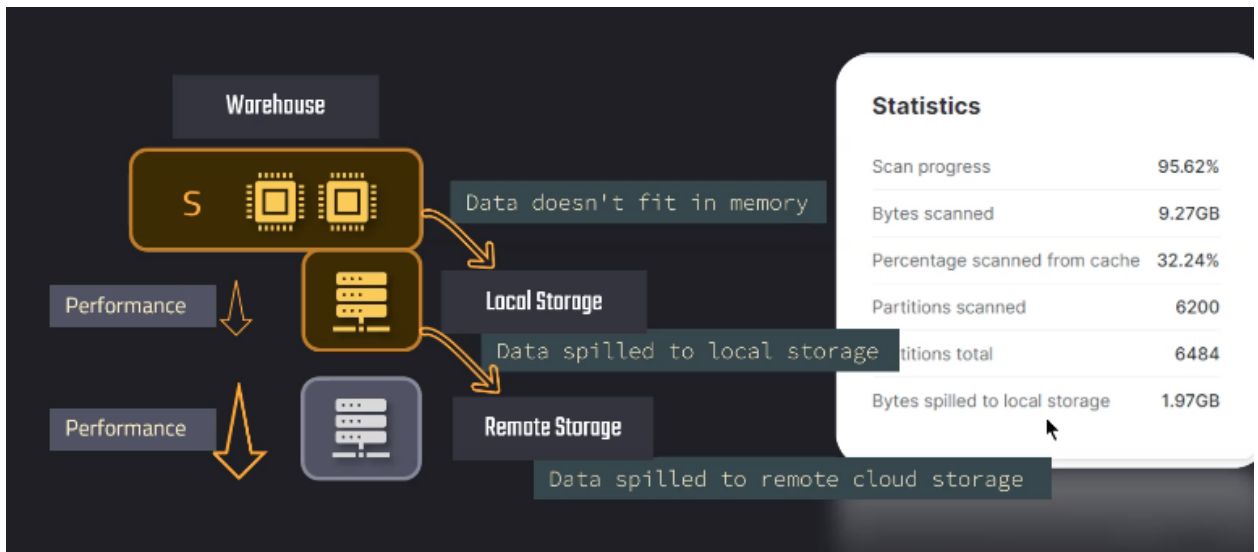
16. And even if our local storage maxes out...then it goes to the remote storage and spills data...

and huge query, then this can be also even spilled to the remote cloud storage.

So then a bucket, for example, if our account is based on an external bucket, will be created and

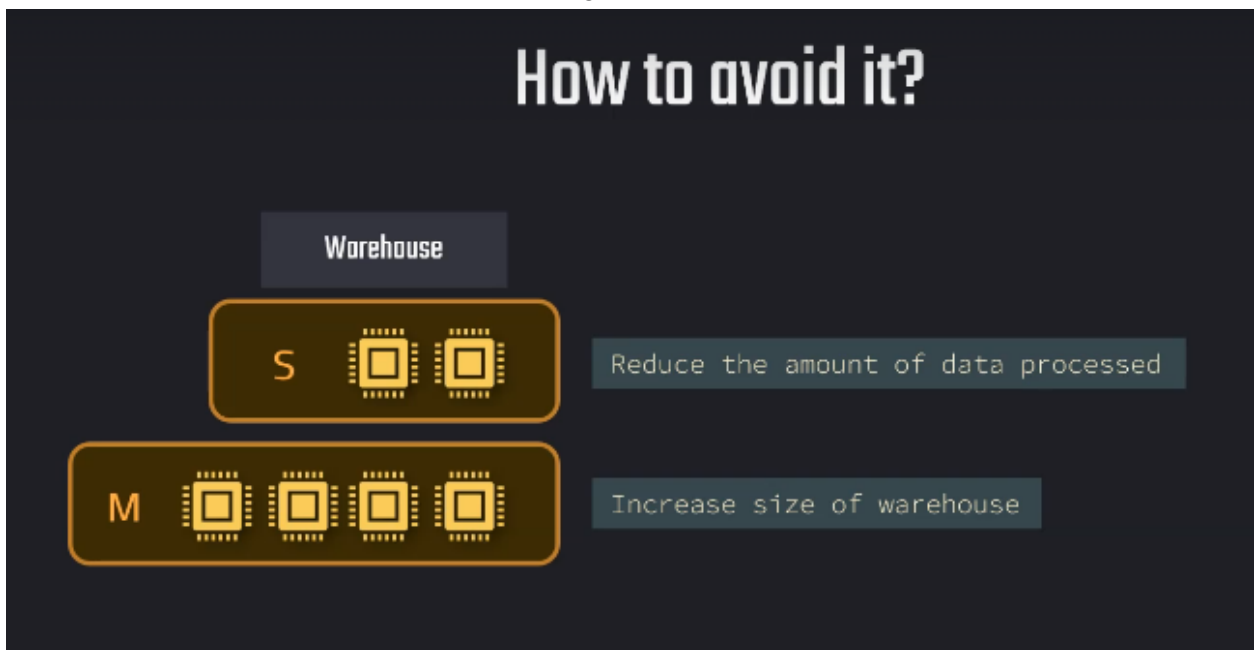
the data will be spilled to this external bucket and this will now even slow down the performance even

more.



17. How can we avoid this data spilling?

18. We can reduce the number of queries running at same time and



19. How to access query history?

Query History in
Snowsight

QUERY_HISTORY
table function in
INFORMATION_SCHEMA

QUERY_HISTORY
view in
ACCOUNT_USAGE schema

Activity

Query History

Copy History

Task History

SELECT percentile_cont(0.4) WITHH

SELECT percentile_cont(0.4) WITHH

SELECT percentile_cont(0.4) WITHH

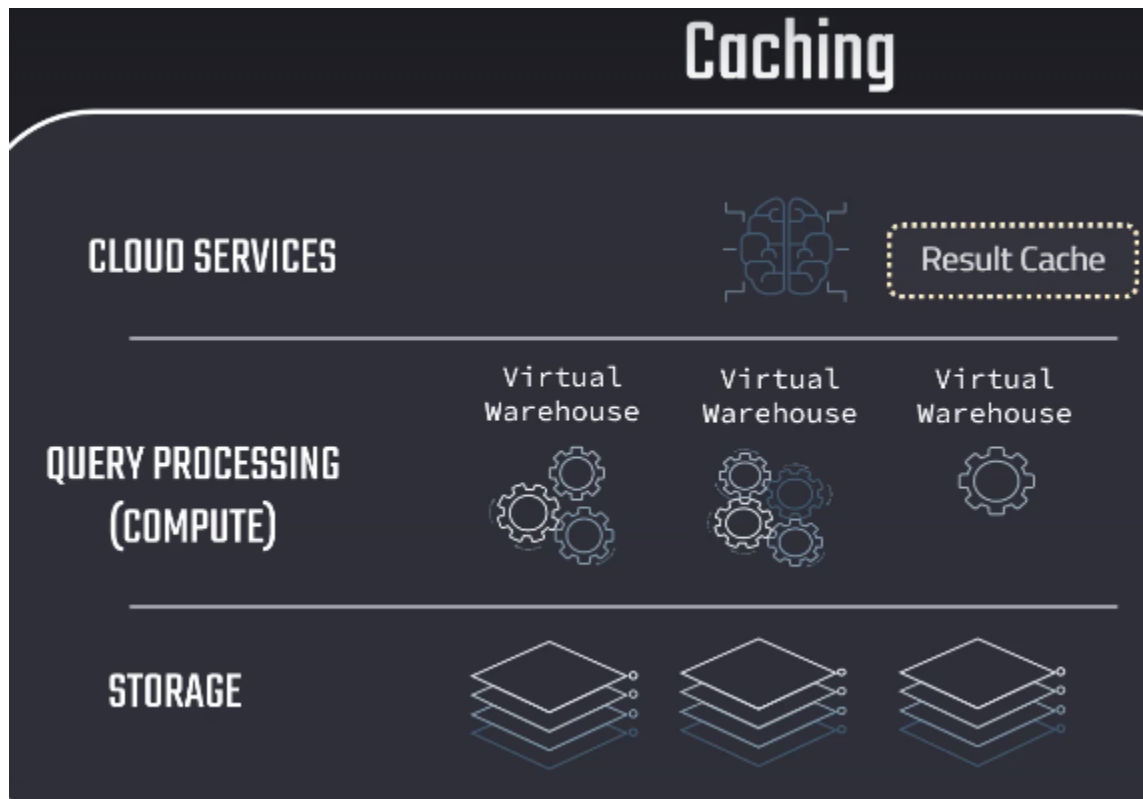
select * from table(information_schema.query_history())
order by start_time;

select * from snowflake.account_usage.query_history


20.

Caching


1. In Snowflake, we have three mechanisms for caching. They are in different locations and we need to understand those three caching mechanisms.



2. Here first we have result cache....it is located in cloud service layer..so we dont need any storage or compute resources for this
3. When we execute a query...a copy of this result..is stored in result cache



4. Next we have data cache.,...also called as virtual warehouse cache..and it is connected directly with our local SSD




So this cannot be

shared with another warehouse.

5. In the query profile, we can see this under the name of local disk I/O.
6. Next we have remote disk...if the data is not cached..then we need to retrieve this data from storage layer..this is not a cache mechanism



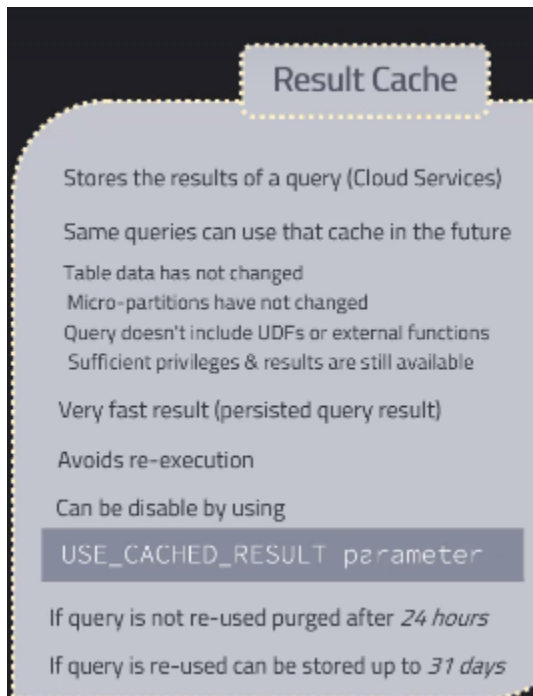
- 7.
8. Then we have metadata cache...it contains statistics for tables, columns and micro



partitions

9. Lets now look into result cache...
10. It stores the result of a query ..when executed...
11. And if the same query runs again..we get the result faster...but the table data must not be modified..
12. The results in result cache are stored for 24hrs....and if run our query again in this 24hrs....then the result will be stored for next 24hrs

13. If query is reused...result cache can store this result for 31days



14. Data Cache

15. It is a local SSD of a given warehouse

16. We can optimize this by using the queries with similar data..in the same warehouse



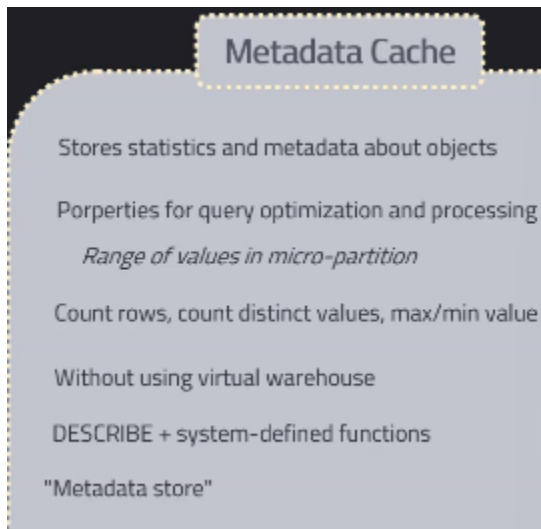
17. Metadata cache

18. It stores the pre calculated statistic of the object...like count of rows, max/min values etc

19. For example it can be range of values in micro-partition

20. And also for commands like DESCRIBE ...this metadata will be helpful

21. This is sometimes referred as a metadata store

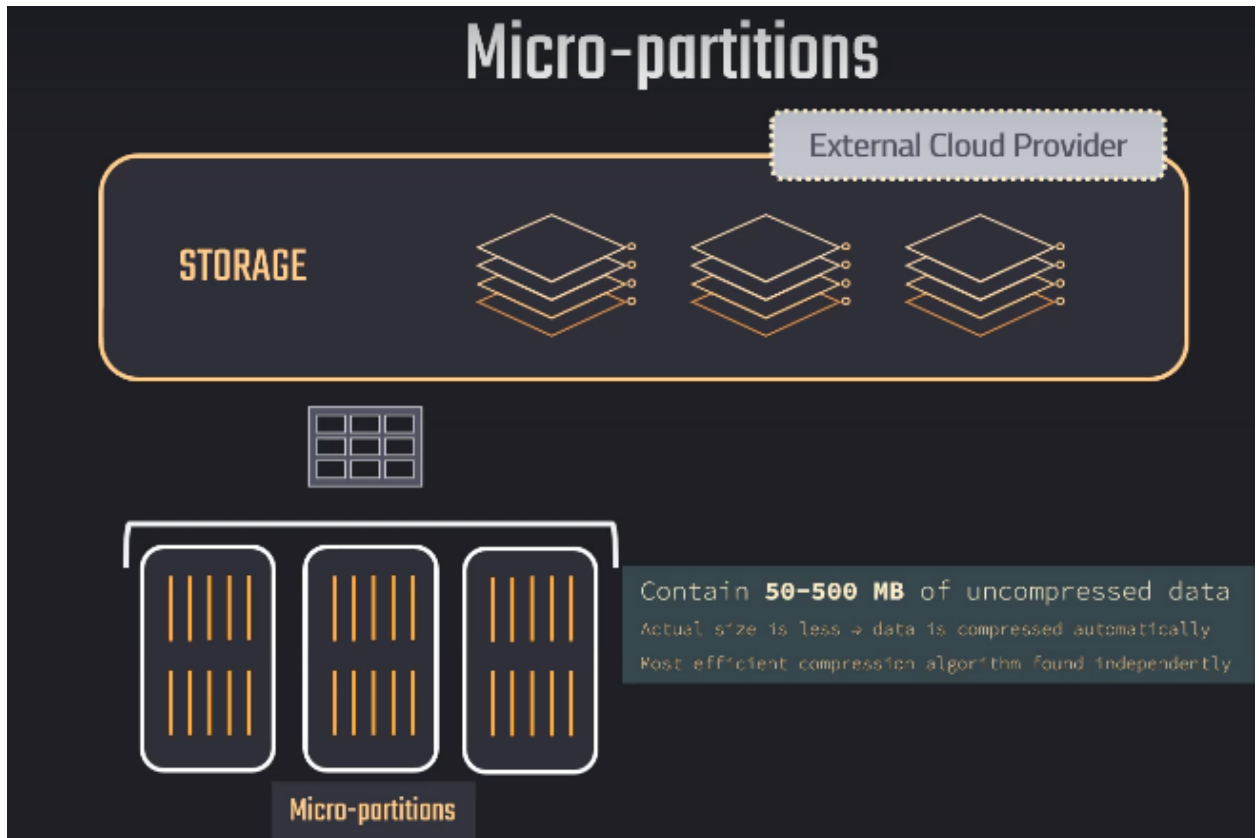


22. And for the virtual private edition, there can be a dedicated metadata store that is not shared with any other customers.

23. For hands on refer file

Micro Partitions

1. Actually our data is stored in the external cloud provider
2. But the question is how is now the data in a table actually stored with this cloud provider?
3. Data is stored in micro-partitions..which is very specific to snowflake architecture

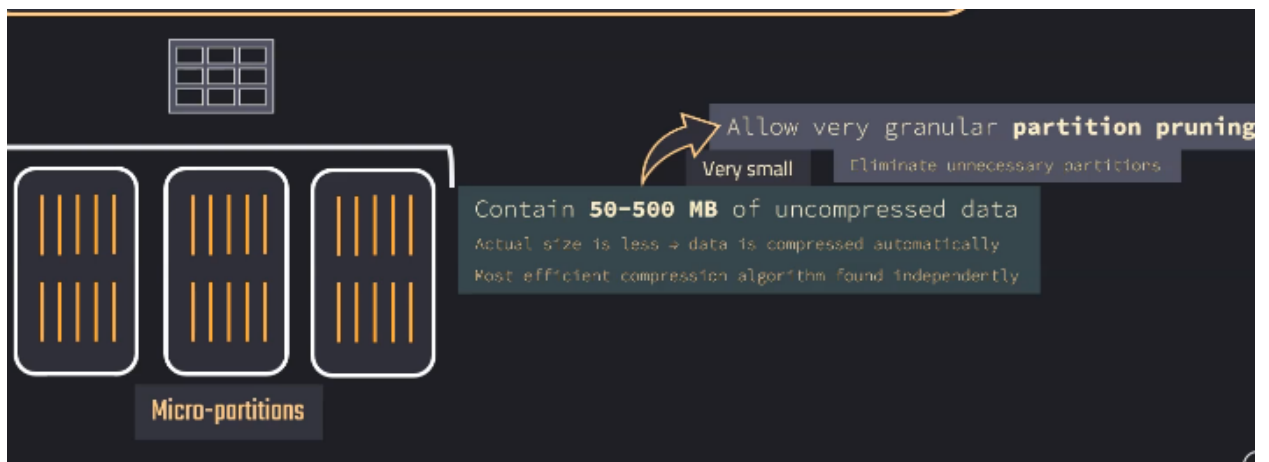


- 4.
5. But actual size of this data is very less...snowflake finds the best compression algo for each partition and compresses each partition
6. Also it allow granular partition pruning...means it eliminates the unnecessary partition required for the query

7. Bard on granular pruning

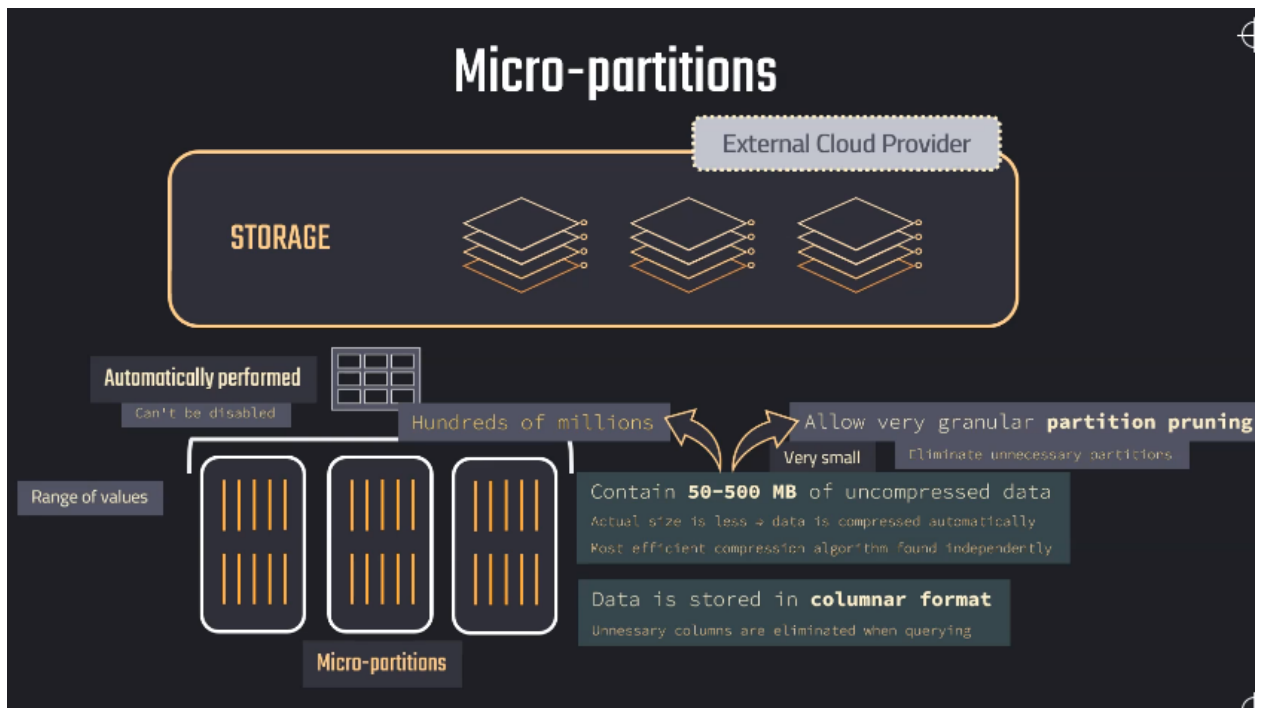
Fine-grained pruning in Snowflake is the process of eliminating micro-partitions from a query scan based on the predicate expressions in the query. Snowflake uses a variety of techniques to perform fine-grained pruning, including:

- **Range pruning:** Snowflake can prune micro-partitions that do not contain data within the specified range of values for a predicate expression. For example, if a query contains a predicate expression such as `order_date BETWEEN '2023-10-21' AND '2023-10-31'`, Snowflake can prune all of the micro-partitions that do not contain data for the specified date range.
- **Equality pruning:** Snowflake can prune micro-partitions that do not contain data equal to the specified value for a predicate expression. For example, if a query contains a predicate expression such as `customer_id = 12345`, Snowflake can prune all of the micro-partitions that do not contain data for the specified customer ID.

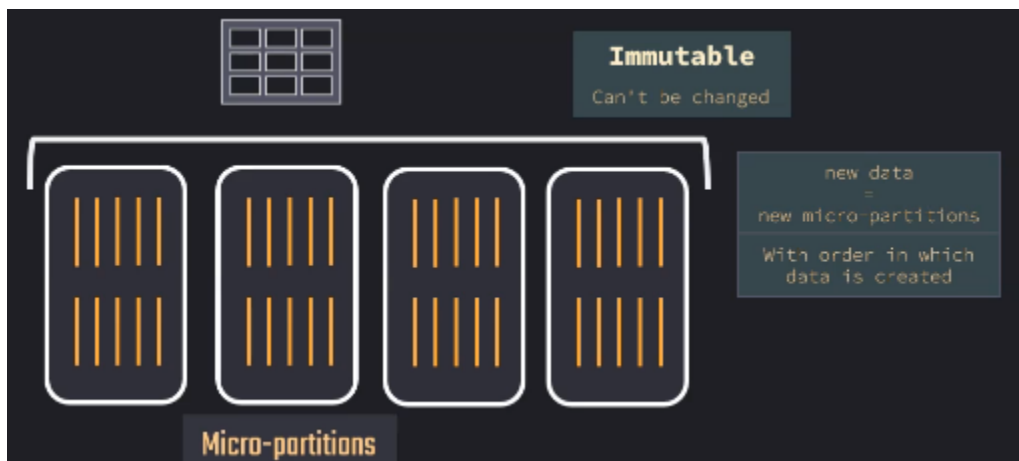


- 8.
9. It may contain hundred of millions of partitions And also the data is stored not in rows but in a so-called columnar format.

10. And this micro-partition is automatically done by snowflake..we cannot disable it



11. These partitions are immutable and..if there's any new data...this will create new micro-partitions and this will happen just with the natural order with which the data is inserted.



12. We can actually locate the data ..using clustering keys

So when we insert new data, it will be just added to a new micro partition and then the next data will

again be added into a new micro partition.

And this is also how we can have an influence on the performance by introducing so-called clustering

keys that can have an influence on how the data is actually located into these micro partitions, because

this can be reorganized by deleting old partitions and inserting the data in a different way into new

micro partitions.

Micro Partition Summary & Benefit

Micro Partition 1		Rows 1 to 4		1	
id	13	9			
	19	4	9		
			4		
					1
f_name	fn-13	fn-9			
	fn-19	fn-4	fn-9		
			fn-4		
					9
l_name	ln-13	ln-9			
	ln-19	ln-4	ln-9		
			ln-4		
					4
dob	2/5/1995	8/13/1997			
	9/11/1997	9/19/1997	/1997		
			/1997		
					9
active	FALSE	TRUE			
	TRUE	TRUE	TRUE		
			TRUE		
					1997
city	San Francisco	Washington			
	Brooklyn	Miami	ington		
					1997
					/1997

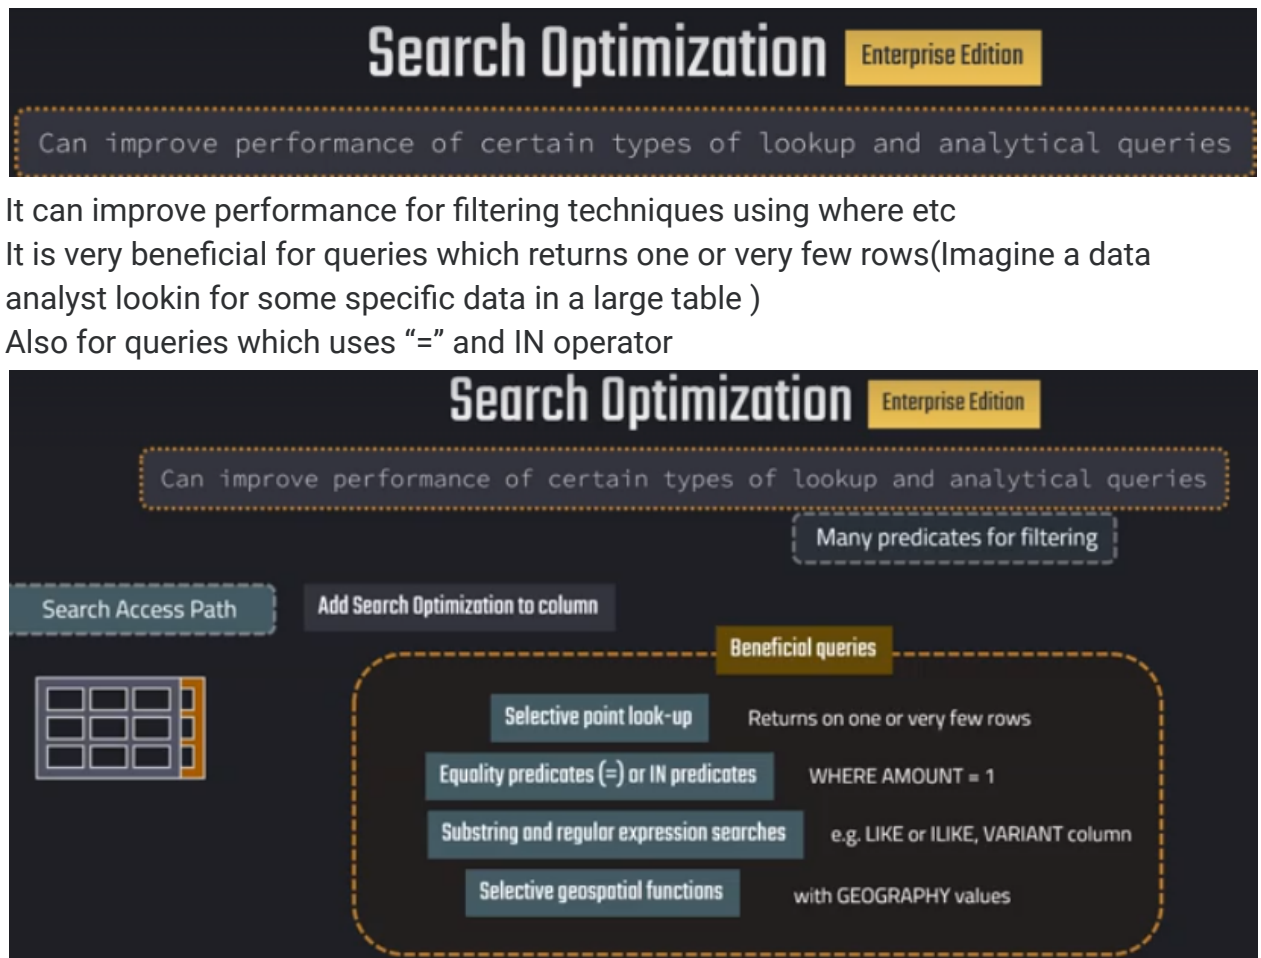
- * Unlike traditional static partitioning, Snowflake **micro-partitions** are **derived automatically**; they don't need to be explicitly defined up-front or maintained by user
- * As the name suggests, micro-partitions are small in size (**50 to 500 MB**, before compression), which enables **extremely efficient DML** and **fine-grained pruning** for faster queries.
- * Micro-partitions can **overlap in their range of values**, which, combined with their uniformly small size, helps prevent skew.
- * Columns are **stored independently** within micro-partitions, often referred to as **columnar storage**. This enables efficient scanning of individual columns; **only the columns referenced by a query** are scanned.
- * Columns are also **compressed individually within micro-partitions**. Snowflake automatically determines the most efficient **compression algorithm** for the columns in each micro-partition.

Imp for exam

For micropartitioning and clustering :

<https://www.youtube.com/watch?v=UNBysn1M9Vg>

Search Optimization Service



5. It is maintained by snowflake..once we enable this..and this also consumes



credits
for maintaining search access

Storage is needed

Add Search Optimization to table

```
ALTER TABLE mytable ADD SEARCH OPTIMIZATION;
```

- 6.
7. We need to have ownership privilege or ADD Search optimization privilege ...to add search optimization to a table

Search optimization in Snowflake is a feature that can improve the performance of point lookup queries. Point lookup queries are queries that return a small number of rows, or even a single row, based on a specific value. Search optimization works by creating a special data structure called a search access path for the columns involved in the point lookup query. The search access path allows Snowflake to quickly locate the micro-partitions containing the necessary data, reducing the amount of time and computing resources required to execute the query.

Search optimization can be enabled for a table or for individual columns within a table. To enable search optimization for a table, you can use the following SQL statement:

SQL

```
ALTER TABLE table_name ADD SEARCH OPTIMIZATION;
```

Use code with caution. [Learn more](#)



- 8.

To enable search optimization for an individual column, you can use the following SQL statement:


SQL

```
ALTER TABLE table_name ADD SEARCH OPTIMIZATION ON column_name;
```

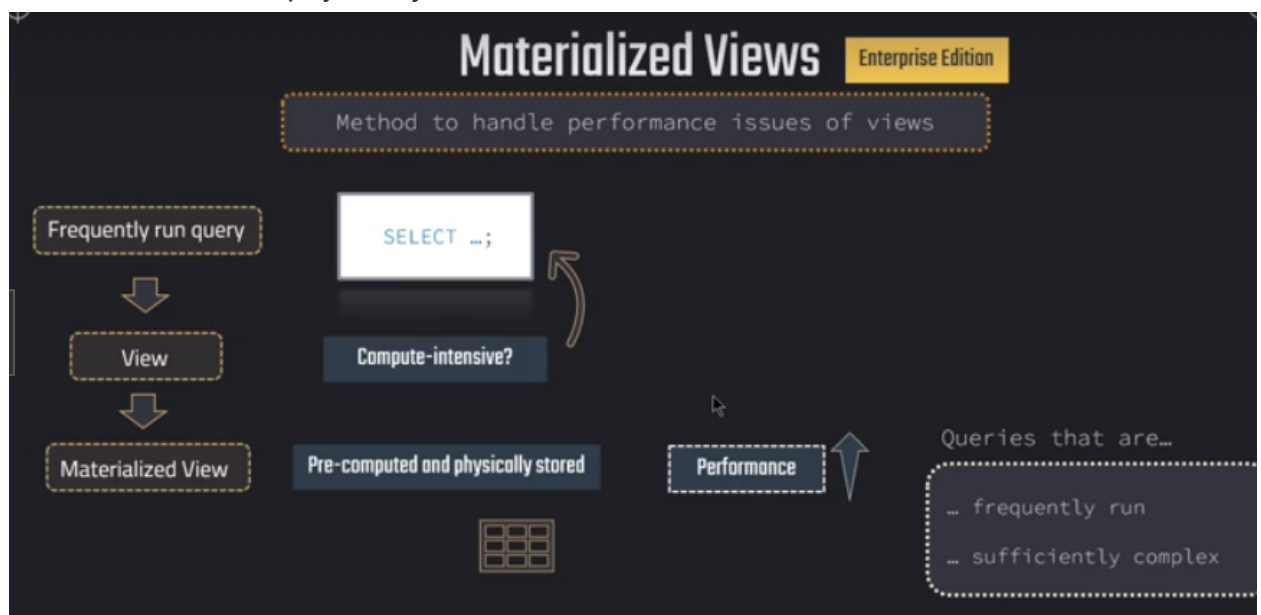
9.

Materialized Views

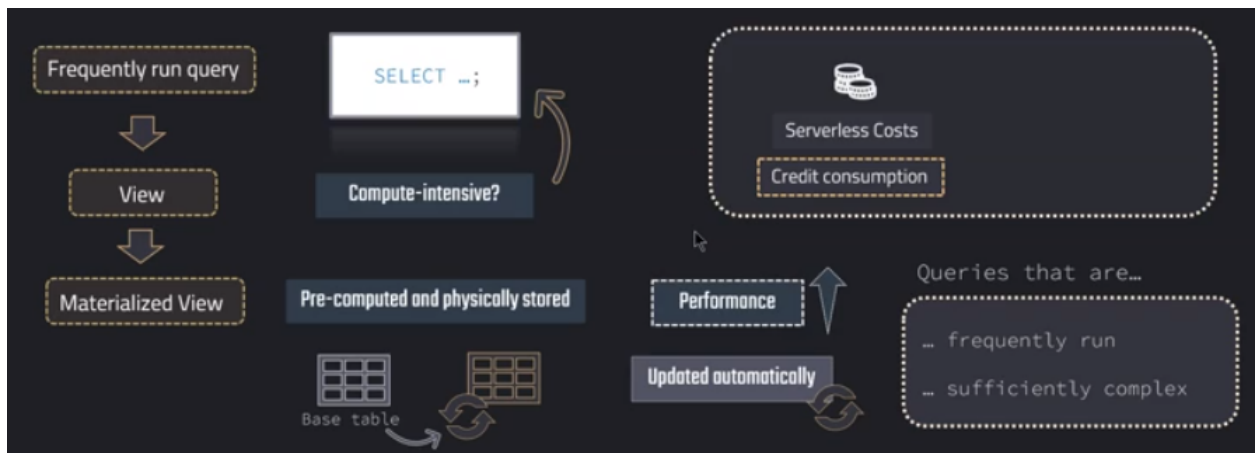
1.



2. Lets assume we have query which is frequently runned to see the data..so we use views
3. But sometimes it takes lot of compute resources...and this causes compute costs and performance
4. The sol to the above prob is Materialized views...which basically precomputes and stores the data physically



5. And if we update our table in the base location...then the data which is in materialized view..will also gets updated automatically...it is done by snowflake server...and we will charged for computing



6. And also we get charged for storage costs..as we are physically storing data now




7. As this materialized views is creating additional costs...we have to be careful and have to start slow..by using one materialized view

8. **Resource monitors can't control Snowflake-managed warehouses**

9. To get addition information we use

```
SELECT * FROM TABLE(INFORMATION_SCHEMA.MATERIALIZED_VIEW_REFRESH_HISTORY());
```

10. Syntax



```
CREATE MATERIALIZED VIEW v_1 AS
SELECT * FROM table1 where c1 = 200
```

- 11.
12. Limitations

13. We cannot use joins while creating MV(mater view)..

14. And cannot create a MV on a view or MV...but we can create MV on external table

Limitations

- Query only 1 table (no joins)
- No views / materialized views
- No window functions, UDFs, HAVING
- Some aggregate functions

Can be created on external tables

15.

16. To decrease some costs we can suspend the MV..or if we dont want this MV..we

```
ALTER MATERIALIZED VIEW v_1 SUSPEND;  
ALTER MATERIALIZED VIEW v_1 RESUME;
```

```
DROP MATERIALIZED VIEW v_1;
```

can drop this

Warehouse Considerations

1. Resizing
2. If we resize(increase compute etc) our warehouse when there are multiple queries running..then the queries which are running will get the old size...only the new queries will get updated warehouse power

Resizing

- Warehouses can be resized even when query is running or when suspended
⇒ Impact only *future* queries, not on the running one

Resizing	<ul style="list-style-type: none">Warehouses can be resized even when query is running or when suspended ⇒ Impact only <i>future</i> queries, not on the running one
Scale up vs. Scale out	<ul style="list-style-type: none">Scale up (resize): More complex queriesScale out: More users (more queries)
Dedicated warehouse	<ul style="list-style-type: none">Isolate workload of specific userDifferent type of workload ⇒ different warehouseEnable auto-suspend and auto-resume (available for all warehouses)

3.